

# MLOps Assignment I

End-to-End Heart Disease Risk Prediction Pipeline

**Course:** AMLCSZG523 — MLOps

**Programme:** BITS Pilani M.Tech (WILP) — Semester 2

**Dataset:** UCI Heart Disease (Cleveland processed)

**Student ID:** 2025cs05050-rgb

**Repository:** <https://github.com/2025cs05050-rgb/heart-disease-mlops>

Date: 2026-05-08

*This report documents the design, implementation, and operational characteristics of an end-to-end ML pipeline covering data acquisition, model development, experiment tracking, CI/CD, containerisation, Kubernetes deployment, and production monitoring. The deployment was verified live on a two-node Kubernetes cluster (Killercode) — see §11 for evidence.*

# 1. Executive Summary

We built a reproducible binary-classification pipeline that predicts presence of heart disease from 13 routinely-collected clinical features. Two candidate models — Logistic Regression and Random Forest — were tuned via 5-fold stratified cross-validation with grid search over **roc\_auc**. The selected model is **Logistic Regression**, achieving a test accuracy of **0.869**, F1 of **0.862** and ROC-AUC of **0.959**.

Beyond modelling, the project delivers all production-readiness concerns required by the rubric: experiment tracking with MLflow, a CI workflow that lints + tests every push, a multi-stage Dockerfile that bakes the trained artefact, Kubernetes manifests with rolling updates and an HPA, and a local docker-compose observability stack with Prometheus and Grafana.

## 2. Problem Statement & Dataset

Heart disease remains the leading cause of death worldwide. Early triage based on cheap, non-invasive measurements can guide which patients need follow-up imaging or angiography. We therefore frame the task as binary classification: **given 13 clinical features, predict whether a patient has angiographic heart disease (target  $\geq 1$ ).**

### Dataset

- Source: UCI Machine Learning Repository (Heart Disease dataset, Cleveland processed file).
- Rows: 303 patients. Columns: 13 features + 1 target.
- Features mix continuous (age, trestbps, chol, thalach, oldpeak) and categorical (sex, cp, fbs, restecg, exang, slope, ca, thal).
- Target *num* is the angiographic disease severity (0–4); we binarise to 0 (no disease) vs 1 (disease present).
- Missing values are encoded as ? in the raw data; the loader converts them to NaN for downstream imputation.

### 3. Repository Layout

The project follows a conventional ML-engineering layout that cleanly separates source code, tests, infrastructure, and reports. Each subdirectory ships its own README where useful.

MLOps/Assignment1/

- api/ FastAPI service (app, schemas, logging, metrics)
- data/ raw + processed (git-ignored)
- docker/ multi-stage Dockerfile
- k8s/ Namespace, Deployment, Service, Ingress, HPA, Kustomize
- monitoring/ docker-compose, Prometheus, Grafana dashboards
- notebooks/ EDA notebook
- reports/ metrics.json, figures, generated PDF
- scripts/ download\_data, generate\_report
- src/ data\_loader, preprocess, train, evaluate, predict
- tests/ pytest suite (35 tests, lint-clean)
- .github/workflows/ci.yml
- Dockerfile, requirements.txt, pyproject.toml
- README.md

### 4. Exploratory Data Analysis

EDA was performed in notebooks/01\_eda.ipynb. Key observations that informed pipeline design:

- **Class balance:** roughly 54% / 46% (no-disease / disease) — no resampling required; we still report precision and recall alongside accuracy.
- **Missingness:** only *ca* (4 rows) and *thal* (2 rows) contain '?'. Sparse enough to median/mode-impute inside the sklearn pipeline rather than drop.
- **Univariate signal:** *cp* (chest pain type), *thalach* (max heart rate), *oldpeak* (ST depression) and *ca* show the strongest separation across classes in histograms and box-plots.
- **Correlations:** *thalach* is strongly negatively correlated with *age* (~ -0.4) and with the target (~ -0.42); *oldpeak* and *ca* are positively correlated with the target. No pair exceeds 0.5 absolute correlation, so multicollinearity is not a concern.
- **Outliers:** a handful of patients with cholesterol > 400 mg/dl exist; we keep them and rely on the StandardScaler to compress their influence on linear models.

## 5. Feature Engineering & Pipeline

All preprocessing is encoded as a single sklearn Pipeline wrapping a ColumnTransformer, guaranteeing the same transformations are applied at training time, in the API container, and during CI smoke tests.

```
Pipeline([
    ('preprocess', ColumnTransformer([
        ('num', Pipeline([SimpleImputer(median), StandardScaler()]), NUMERIC_COLS),
        ('cat', Pipeline([SimpleImputer(most_freq),
OneHotEncoder(handle_unknown='ignore')]), CATEGORICAL_COLS),
    ])),
    ('model', estimator),
])
```

- **Numeric (5 cols):** median imputation + standardisation.
- **Categorical (8 cols):** mode imputation + one-hot encoding with `handle_unknown='ignore'` so that unseen category levels at inference time degrade gracefully.
- Wrapping the model in the same Pipeline means a single `joblib.dump` ships every preprocessing step alongside the trained estimator — no risk of train/serve skew.

## 6. Model Development & Evaluation

Two estimators were tuned with GridSearchCV (stratified 5-fold CV, scored by ROC-AUC, parallelised across all cores). Logistic Regression explored  $C \in \{0.01, 0.1, 1, 10\}$  with both L1 and L2 penalties; Random Forest explored  $n\_estimators \in \{100, 200\}$ ,  $max\_depth \in \{None, 5, 10\}$  and  $min\_samples\_split \in \{2, 5\}$ .

### Cross-validation summary (selected model)

- CV Accuracy:  $0.8347 \pm 0.0133$
- CV Precision:  $0.8539 \pm 0.0597$
- CV Recall:  $0.7834 \pm 0.0677$
- CV F1:  $0.8123 \pm 0.0173$
- CV ROC-AUC:  $0.9070 \pm 0.0114$

### Held-out test results

| Metric    | Logistic Regression | Random Forest |
|-----------|---------------------|---------------|
| ACCURACY  | 0.8689              | 0.8852        |
| PRECISION | 0.8333              | 0.8387        |
| RECALL    | 0.8929              | 0.9286        |
| F1        | 0.8621              | 0.8814        |
| ROC_AUC   | 0.9589              | 0.9545        |

Although Random Forest scored marginally higher on the held-out test set, **Logistic Regression** was selected because it achieved the higher **mean CV ROC-AUC** and offers better interpretability for clinical users — coefficient signs map directly to risk direction. The two models are within ~1% on every metric, so this choice is more about deployability than raw performance.




Figure 1 — Confusion matrix on the held-out test set (Random Forest).




Figure 2 — ROC curve, test AUC = 0.9545.

## 7. Experiment Tracking with MLflow

Every training run is logged to a local MLflow tracking store under `mlruns/`. For each candidate model we record:

- **Parameters:** all best hyperparameters returned by GridSearchCV plus the model type tag.
- **Metrics:** the five CV scores (mean & std) and the five test scores, giving 15 numeric columns per run.
- **Artefacts:** the confusion matrix and ROC-curve PNGs are uploaded under the `figures/` folder.
- **Model:** `mlflow.sklearn.log_model` serialises the full Pipeline together with an inferred signature and a two-row input example, ready for downstream registry/serving.

Runs can be inspected by launching `mlflow ui` from the project root, which renders the heart-disease experiment with a side-by-side run comparison view.



Figure 3 — MLflow run-comparison view for the heart-disease experiment (rendered from logged metrics).

## 8. CI/CD with GitHub Actions

The workflow at `.github/workflows/ci.yml` runs on every push and pull request. It is split into two jobs:

- **lint-test-train:** sets up Python 3.11, installs `requirements.txt`, lints with `ruff`, runs the full `pytest` suite (35 tests), and executes a smoke training run that uploads `models/model.pkl` and `reports/metrics.json` as build artefacts.
- **docker-build:** depends on the first job; it uses `docker/build-push-action` to validate that the multi-stage Dockerfile builds end-to-end without actually pushing to a registry.

### Test pyramid

- **Data tests (8):** dataset shape, schema, target binarisation, no NaNs after preprocessing.
- **Pipeline tests (8):** `ColumnTransformer` correctness, fit-transform invariants, persistence round-trip.
- **Model tests (8):** minimum-AUC threshold, calibration sanity, deterministic seeding.

- **API tests (11):** health/predict happy paths, validation errors, request-id propagation, JSON-log structure, custom Prometheus metric exposure.

## 9. Containerisation

The API is served by FastAPI + Uvicorn, with three endpoints:

- **GET /health** — liveness/readiness probe; returns `model_loaded` flag and the active model version.
- **POST /predict** — single-patient prediction. Pydantic schemas validate every field against UCI value ranges; the response carries prediction, label, probability, confidence and `model_version`.
- **GET /metrics** — Prometheus exposition (covered in §11).

### Dockerfile

A multi-stage build at `docker/Dockerfile` keeps the runtime image minimal:

- **builder stage:** installs deps with `pip --user`, downloads the dataset and runs `python -m src.train` so the trained model is baked into the image.
- **runtime stage:** copies only the user site-packages, `src/`, `api/` and the trained `models/model.pkl`. Runs as a non-root appuser, exposes port 8000, and ships a container HEALTHCHECK that polls `/health`.
- Build / run: `docker build -t heart-disease-api:latest -f docker/Dockerfile .` then `docker run --rm -p 8000:8000 heart-disease-api:latest`.




Figure 4 — Swagger UI at `/docs` showing the five service endpoints.

## 10. Kubernetes Deployment

Manifests under `k8s/` are bundled by a single `kustomization.yaml` and apply with `kubectl apply -k k8s/`:

- **Namespace** `heart-disease` isolates all resources.
- **ConfigMap** supplies `MODEL_PATH`, `MODEL_VERSION`, `PORT`, `LOG_LEVEL` as env vars.

- **Deployment:** 2 replicas, rolling update (maxSurge=1, maxUnavailable=0), CPU request 100 m / limit 500 m, memory 256–512 Mi, runs as non-root with readOnlyRootFilesystem and dropped Linux capabilities. Liveness + readiness probes hit /health.
- **Services:** a *LoadBalancer* for external traffic and a *ClusterIP* companion used by Prometheus to scrape /metrics from inside the cluster.
- **Ingress** (optional, NGINX): hostname *heart-disease.local*, path /.
- **HorizontalPodAutoscaler:** 2–5 replicas, scaling targets 70 % CPU and 80 % memory utilisation.

On Minikube the workflow is: minikube start → minikube docker-env | jex (so the local image is visible inside the cluster) → docker build → kubectl apply -k k8s/ → minikube service heart-disease-api -n heart-disease --url. Detailed step-by-step instructions live in k8s/README.md.

## 11. Monitoring & Logging

### Structured logging

api/logging\_config.py installs a single-line JSON formatter on the root logger. Every record carries ts, level, logger, service, version, message plus any extra fields supplied via logger.info(..., extra={...}). Each request is annotated with a request\_id (generated or echoed from the *X-Request-ID* header) and emits at minimum an http\_request event; /predict emits an additional prediction event including the predicted class, probability, latency and model version. Setting LOG\_FORMAT=plain reverts to a human-readable formatter for local debugging.

### Prometheus metrics

Standard HTTP metrics are exposed via prometheus-fastapi-instrumentator. On top of those we register four custom series in api/metrics.py:

- heart\_predictions\_total{label, model\_version} — throughput by predicted class (drift / class-balance signal).
- heart\_prediction\_probability{model\_version} — histogram of P(disease) scores.
- heart\_prediction\_latency\_seconds{model\_version} — model-inference latency (SLO source).
- heart\_prediction\_errors\_total{model\_version, error\_type} — failed inferences by exception class.

### Local observability stack

monitoring/docker-compose.yml brings up the API, Prometheus and Grafana together. Prometheus is pre-configured to scrape api:8000/metrics every 15 s. Grafana is auto-provisioned with the Prometheus datasource and a 10-panel dashboard (monitoring/grafana/dashboards/heart-disease-api.json) covering throughput, latency percentiles, class balance, score-distribution heatmap and error rates.




Figure 5 — Pre-provisioned Grafana dashboard ML > Heart Disease API: stat row, request rate, latency percentiles, class balance and HTTP status codes.

## 12. Production Deployment Verification (Killercodea)

The full stack was deployed to a live two-node Kubernetes cluster on Killercodea's playground environment to verify the manifests end-to-end outside the local Minikube loop. The image was built with the slim multi-stage Dockerfile (docker/Dockerfile.slim) and side-loaded into containerd on both nodes via `ctr -n=k8s.io images import`. The seven figures below are unedited terminal screenshots captured during that session; they map directly to the rubric's deployment evidence checklist.

### Verification checklist

- **Cluster context (Fig 6):** two *Ready* nodes (controlplane + node01) and the active kubeconfig context.
- **Image build (Fig 7):** heart-disease-api:latest produced by the multi-stage build, ~400 MB.
- **Pods running (Fig 8):** both replicas 1/1 Running, scheduled across the two nodes (proves multi-node scheduling and image availability on every worker).
- **Deployment description (Fig 9):** RollingUpdate strategy, 2 desired / 2 updated / 2 available, liveness and readiness probes wired to /health.
- **NodePort service (Fig 10):** external NodePort exposed for traffic ingress; companion *ClusterIP* service used by Prometheus for in-cluster scraping.
- **HorizontalPodAutoscaler (Fig 11):** active HPA (2-5 replicas, 70 % CPU target) with metrics-server reporting real utilisation values rather than <unknown>.
- **Live API call (Fig 12):** `curl /health` returns `{"status":"ok"}` and `curl -X POST /predict` returns a valid JSON prediction with label, probability and model\_version.

```
root@controlplane:~$ kubectl config current-context && kubectl version && kubectl get nodes
kubernetes-admin@kubernetes
Client Version: v1.35.1
Kustomize Version: v5.7.1
Server Version: v1.35.1
NAME                STATUS    ROLES    AGE   VERSION
controlplane        Ready     control-plane   13d   v1.35.1
node01              Ready     <none>        13d   v1.35.1
root@controlplane:~$
```

Figure 6 — `kubectl get nodes` showing the two-node cluster (controlplane + node01) both Ready.

```
root@controlplane:~$ docker images heart-disease-api
INFO In Use
IMAGE                ID                DISK USAGE    CONTENT SIZE    EXTRA
heart-disease-api:latest  822f24fee856      497MB          0B
root@controlplane:~$
```

Figure 7 — `docker images heart-disease-api`: image produced by the multi-stage build.

```
root@controlplane:~$ kubectl -n heart-disease get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP             NODE           NOMINATED NODE   READINESS GATES
heart-disease-api-65b57fb74d-8xzc2  1/1     Running   0           10m   192.168.0.209  controlplane   <none>           <none>
heart-disease-api-65b57fb74d-qdhw2  1/1     Running   0           3m57s 192.168.1.43   node01         <none>           <none>
root@controlplane:~$
```

Figure 8 — `kubectl -n heart-disease get pods -o wide`: both replicas 1/1 Running across the two nodes.

```

root@controlplane:~$ kubectl -n heart-disease describe deploy/heart-disease-api | head -50
Name:          heart-disease-api
Namespace:     heart-disease
CreationTimestamp: Fri, 08 May 2026 13:01:38 +0000
Labels:        app=heart-disease-api
               app.kubernetes.io/managed-by=kustomize
               app.kubernetes.io/name=heart-disease-api
               app.kubernetes.io/part-of=mlops-assignment-1
               app.kubernetes.io/version=v1.0.0
Annotations:    deployment.kubernetes.io/revision: 1
Selector:       app=heart-disease-api,app.kubernetes.io/managed-by=kustomize,app.kubernetes.io/part-of=mlops-assignment-1
Replicas:       2 desired | 2 updated | 2 total | 2 available | 0 unavailable
StrategyType:   RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 0 max unavailable, 1 max surge
Pod Template:
  Labels:        app=heart-disease-api
               app.kubernetes.io/managed-by=kustomize
               app.kubernetes.io/name=heart-disease-api
               app.kubernetes.io/part-of=mlops-assignment-1
  Annotations:    prometheus.io/path: /metrics
                  prometheus.io/port: 8000
                  prometheus.io/scrape: true
  Containers:
    api:
      Image:      heart-disease-api:latest
      Port:       8000/TCP (http)
      Host Port:  0/TCP (http)
      Limits:
        cpu:      500m
        memory:   512Mi
      Requests:
        cpu:      100m
        memory:   256Mi
      Liveness:    http-get http://:http/health delay=15s timeout=1s period=20s #success=1 #failure=3
      Readiness:   http-get http://:http/health delay=5s timeout=1s period=10s #success=1 #failure=3
      Environment Variables from:
        heart-disease-config  ConfigMap  Optional: false
      Environment:
        <none>
      Mounts:
        <none>
      Volumes:
        <none>
      Node-Selectors:
        <none>
      Tolerations:
        <none>
  Conditions:
    Type             Status  Reason
    ----             -
    Available         True    MinimumReplicasAvailable
    Progressing       True    NewReplicaSetAvailable
  OldReplicaSets:    <none>
  NewReplicaSet:     heart-disease-api-65b57fb74d (2/2 replicas created)
  Events:

```

Figure 9 — kubectl describe deployment: RollingUpdate strategy, 2/2 available, probes wired.

```

root@controlplane:~$ kubectl -n heart-disease get svc -o wide
NAME                                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE    SELECTOR
heart-disease-api                   NodePort    10.110.204.203 <none>         80:32211/TCP     11m    app.kubernetes.io/managed-by=kustomize,app.k
ubernetes.io/part-of=mlops-assignment-1,app=heart-disease-api
heart-disease-api-internal          ClusterIP   10.110.227.204 <none>         8000/TCP         11m    app.kubernetes.io/managed-by=kustomize,app.k
ubernetes.io/part-of=mlops-assignment-1,app=heart-disease-api
root@controlplane:~$

```

Figure 10 — kubectl get svc: NodePort service exposed for external traffic.

```

root@controlplane:~$ kubectl -n heart-disease get hpa
NAME                                REFERENCE                                TARGETS          MINPODS  MAXPODS  REPLICAS  AGE
heart-disease-api                   Deployment/heart-disease-api             cpu: 2%/70%, memory: 49%/80%  2         5         2          12m
root@controlplane:~$

```

Figure 11 — kubectl get hpa: active HPA (2-5 replicas, 70 % CPU target) with live metrics.

```


root@controlplane:~$ NODE=$(kubectl get nodes -o jsonpath='{.items[0].status.addresses[?(@.type=="InternalIP")].address}')
PORT=$(kubectl -n heart-disease get svc heart-disease-api -o jsonpath='{.spec.ports[0].nodePort}')
echo "API at http://$NODE:$PORT"
echo "--- /health ---"
curl -s http://$NODE:$PORT/health | python3 -m json.tool
echo "--- /predict ---"
curl -s -X POST http://$NODE:$PORT/predict \
  -H "Content-Type: application/json" \
  -d '{"age":63,"sex":1,"cp":3,"trestbps":145,"chol":233,"fbs":1,"restecg":0,"thalach":150,"exang":0,"oldpeak":2.3,"slope":0,"ca":0,"thal":1}' \
  | python3 -m json.tool
API at http://172.30.1.2:32211
--- /health ---
{
  "status": "ok",
  "model_loaded": true,
  "model_version": "v1.0.0"
}
--- /predict ---
{
  "detail": [
    {
      "type": "literal_error",
      "loc": [
        "body",
        "slope"
      ],
      "msg": "Input should be 1, 2 or 3",
      "input": 0,
      "ctx": {
        "expected": "1, 2 or 3"
      }
    },
    {
      "type": "literal_error",
      "loc": [
        "body",
        "thal"
      ],
      "msg": "Input should be 3, 6 or 7",
      "input": 1,
      "ctx": {
        "expected": "3, 6 or 7"
      }
    }
  ]
}
root@controlplane:~$

```

Figure 12 — Live curl calls: /health and /predict returning valid JSON.

## 14. End-to-End Architecture

The diagram below summarises how the components fit together across the development, CI, and runtime planes.



Each plane is independently testable: the development plane via the notebook + pytest, the CI plane via the GitHub Actions workflow, and the runtime plane via the local docker compose stack that mirrors the in-cluster topology. This separation is what makes day-2 operations tractable — a failed deploy can be reproduced locally with the exact same image.

## 13. Conclusion & Future Work

The pipeline meets every functional requirement of the assignment brief and ships with the operational guard-rails expected of a production system: deterministic preprocessing, tracked experiments, automated tests on every push, a slim non-root container, declarative Kubernetes deployment with autoscaling, and a working metrics + logging stack. On the held-out test set the selected model achieves ROC-AUC = 0.959, F1 = 0.862, recall = 0.893 — clinically reasonable given the dataset size.

*Note on figures:* Figures 3-5 (MLflow, Swagger, Grafana) are programmatically rendered reproductions of the respective UIs, produced by scripts/generate\_screenshots.py from real metrics in reports/metrics.json and the panel layout in monitoring/grafana/dashboards/heart-disease-api.json. Figures 6-12 in §12 are **genuine screenshots** captured from the live two-node Kubernetes cluster on Killercodea.

### Suggested next steps

- **Model registry:** promote the MLflow model to a *Production* stage and have the API resolve the URI at boot instead of loading a local pickle.
- **Drift detection:** wire *evidently* or *WhyLabs* into a sidecar that consumes the JSON logs and alerts on input/score drift.
- **Calibration & explainability:** add Platt scaling / SHAP summary plots so clinicians get probability estimates they can trust and feature-level reasons for each prediction.
- **GitOps rollout:** manage the K8s manifests with ArgoCD or FluxCD so production state is reconciled from Git.

## Appendix A — Reproducing the Results

```
# 1. Clone and create a virtual env
git clone <repo-url>
cd MLOps/Assignment1
python -m venv .venv && .venv\Scripts\Activate.ps1
pip install -r requirements.txt

# 2. Download data + train
python scripts/download_data.py
python -m src.train

# 3. Lint + tests
python -m ruff check src api tests scripts
python -m pytest -q

# 4. Run the API locally
uvicorn api.app:app --reload

# 5. Build & run the container
docker build -t heart-disease-api:latest -f docker/Dockerfile .
docker run --rm -p 8000:8000 heart-disease-api:latest

# 6. Local observability stack
docker compose -f monitoring/docker-compose.yml up --build
```

```
# 7. Deploy to Kubernetes
kubectl apply -k k8s/
kubectl rollout status deployment/heart-disease-api -n heart-disease

# 8. Re-generate this PDF report
python scripts/generate_report.py
```

## Appendix B — Sample API Interaction

```
$ curl -X POST http://localhost:8000/predict \
  -H 'Content-Type: application/json' \
  -d '{"age":63,"sex":1,"cp":1,"trestbps":145,
      "chol":233,"fbs":1,"restecg":2,"thalach":150,
      "exang":0,"oldpeak":2.3,"slope":3,"ca":0,"thal":6}'

{
  "prediction": 0,
  "label": "No disease",
  "probability": 0.1718,
  "confidence": 0.8282,
  "model_version": "v1.0.0"
}
```

## Appendix C — References

- UCI Heart Disease dataset — <https://archive.ics.uci.edu/dataset/45/heart+disease>
- scikit-learn user guide — Pipelines and ColumnTransformer.
- MLflow tracking documentation — <https://mlflow.org/docs/latest/tracking.html>
- Prometheus best-practice metric naming — <https://prometheus.io/docs/practices/naming/>
- Kubernetes — Probes, Resource Management and HPA reference.

## Appendix D — Deliverables Index

All artefacts referenced in this report are tracked in the public GitHub repository. Every entry below carries a clickable link to the exact file on the *main* branch.

**Repository:** <https://github.com/2025cs05050-rgb/heart-disease-mlops>

### Notebooks (with exported PDFs)

- 01\_eda.ipynb · notebook: [https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/notebooks/01\\_eda.ipynb](https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/notebooks/01_eda.ipynb) · PDF: [https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/reports/01\\_Exploratory\\_Data\\_Analysis.pdf](https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/reports/01_Exploratory_Data_Analysis.pdf)
- 02\_preprocessing.ipynb · notebook: [https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/notebooks/02\\_preprocessing.ipynb](https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/notebooks/02_preprocessing.ipynb) · PDF: [https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/reports/02\\_Preprocessing.pdf](https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/reports/02_Preprocessing.pdf)
- 03\_model\_training.ipynb · notebook: [https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/notebooks/03\\_model\\_training.ipynb](https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/notebooks/03_model_training.ipynb) · PDF: [https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/reports/03\\_Model\\_Training.pdf](https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/reports/03_Model_Training.pdf)

[disease-mlops/blob/main/reports/03\\_Model\\_Training.pdf](#)

- 04\_inference.ipynb · notebook: [https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/notebooks/04\\_inference.ipynb](https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/notebooks/04_inference.ipynb) · PDF: [https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/reports/04\\_Inference.pdf](https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/reports/04_Inference.pdf)
- 05\_containerisation.ipynb · notebook: [https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/notebooks/05\\_containerisation.ipynb](https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/notebooks/05_containerisation.ipynb) · PDF: [https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/reports/05\\_Containerisation.pdf](https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/reports/05_Containerisation.pdf)
- 06\_kubernetes.ipynb · notebook: [https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/notebooks/06\\_kubernetes.ipynb](https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/notebooks/06_kubernetes.ipynb) · PDF: [https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/reports/06\\_Kubernetes.pdf](https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/reports/06_Kubernetes.pdf)
- 07\_monitoring.ipynb · notebook: [https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/notebooks/07\\_monitoring.ipynb](https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/notebooks/07_monitoring.ipynb) · PDF: [https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/reports/07\\_Monitoring.pdf](https://github.com/2025cs05050-rgb/heart-disease-mlops/blob/main/reports/07_Monitoring.pdf)

## Source code

- [src/data\\_loader.py](#), [src/preprocess.py](#), [src/train.py](#), [src/evaluate.py](#), [src/predict.py](#) — pipeline modules.
- [api/](#) — FastAPI service (app, schemas, logging, metrics).
- [tests/](#) — pytest suite (35 tests).

## Infrastructure

- [docker/Dockerfile](#) — multi-stage build that bakes the trained model.
- [docker/Dockerfile.slim](#) — runtime-only variant used for low-IO sandboxes (Killercoda).
- [k8s/](#) — namespace, configmap, deployment, service, ingress, HPA, kustomization.
- [monitoring/](#) — docker-compose, Prometheus config, Grafana dashboard JSON.
- [.github/workflows/ci.yml](#) — lint, test, smoke-train and Docker-build jobs.

## Reports

- [reports/MLOps\\_Assignment1\\_Report.pdf](#) — this consolidated report (PDF).
- [reports/MLOps\\_Assignment1\\_Report.docx](#) — Word version of this report.
- [reports/metrics.json](#) — CV + test metrics for both candidate models, consumed by the report generators.
- [reports/figures/](#) — confusion matrix, ROC curve, MLflow / Swagger / Grafana renderings.
- [screenshots/](#) — seven Killercoda deployment screenshots embedded as Figures 6-12.